

Rohan Chhetry

Sydney, NSW | 0435 265 565 | chhetry.rohan5@gmail.com | [LinkedIn](#) | [Portfolio](#)

HIGHLIGHTS

Data Engineer with 4+ years of experience designing and delivering scalable cloud-native data platforms across Azure, Google Cloud Platform (GCP) and AWS. Experienced in building enterprise ETL/ELT pipelines using Azure Data Factory, Databricks, PySpark and SQL, with hands-on expertise in Infrastructure as Code (Terraform), GitLab CI/CD and cloud-native data integration. Skilled in developing reliable, production-grade data platforms supporting analytics, governance and AI workloads.

- Built enterprise cloud-native data ingestion platforms using Azure Data Factory, Databricks, PySpark and Google Cloud services.
- Automated cloud infrastructure deployment through Terraform and GitLab CI/CD following Infrastructure as Code (IaC) practices.
- Developed scalable data engineering solutions supporting analytics, governance, AI and machine learning workloads across Azure, GCP and AWS.

TECHNICAL SKILLS:

Languages: Python, SQL, JavaScript

Data Engineering: Azure Data Factory (ADF), Databricks, PySpark, Apache Spark, ETL/ELT, Incremental Data Processing, Data Integration

Data Architecture: Lakehouse Architecture, Medallion Architecture (Bronze/Silver/Gold), Data Lakes, Distributed Data Processing

Cloud Platforms: Microsoft Azure, Google Cloud Platform (BigQuery, Pub/Sub, Cloud Storage, Cloud Run, DLP), AWS (Lambda, S3, Glue, Bedrock)

Databases: BigQuery, Snowflake, PostgreSQL

Infrastructure & DevOps: Terraform, GitLab CI/CD, Git, Docker, Azure DevOps

EXPERIENCE

INFOSYS (Client: OPTUS)

SYDNEY, NSW

Data Engineer

April 2026 – Present

- Developed enterprise data ingestion pipelines using Azure Data Factory (ADF), Databricks, PySpark and SQL within an enterprise Medallion (Bronze/Silver/Gold) data platform.
- Built cloud-native ingestion pipelines integrating Google Cloud Pub/Sub, BigQuery and Cloud Storage with Azure Data Factory and Databricks for enterprise analytics workloads.
- Provisioned and managed Google Cloud infrastructure using Terraform and Infrastructure as Code (IaC), with deployments automated through GitLab CI/CD pipelines.
- Implemented Google Cloud DLP masking to identify and de-identify PII, ensuring compliance with enterprise data governance standards.
- Developed reusable parameter-driven ETL workflows and Databricks notebooks using the Optus enterprise data engineering framework.
- Designed and maintained staging simulator environments to validate new ingestion features and framework enhancements before production deployment.
- Implemented logging, monitoring and operational reporting to improve pipeline reliability and observability.

UniPhi Software

SYDNEY, NSW

Software Engineer | Data

July 2025 – Nov 2025

- Designed and implemented Azure Data Factory (ADF) pipelines for orchestrating end-to-end data workflows including ingestion, transformation, and ML retraining.
- Built PySpark transformation pipelines on a Spark/Databricks-style environment, implementing Bronze, Silver, and Gold lakehouse layers and writing optimized Parquet datasets for downstream analytics and machine learning workloads.
- Implemented incremental ingestion pipelines using partition-based processing to efficiently capture new and updated records from source systems.
- Integrated Azure ML workflows to automate model retraining and deployment.
- Implemented pipeline logging and monitoring using Azure-native observability tools to improve reliability and troubleshooting.
- Developed reusable modular ETL components enabling faster development and maintainability across pipelines.

Genese Solution

LONDON, UNITED KINGDOM

Senior Software Engineer | Data

April 2024 – June 2025

- Architected a serverless microservices architecture using AWS Lambda, S3, and Glue to handle scalable event-driven workloads.
- Built Python-based APIs enabling natural language queries that translate into validated SQL queries.
- Integrated Amazon Bedrock LLMs with a Pinecone vector database for semantic retrieval, enabling natural-language-to-SQL generation.
- Implemented validation layers to sanitize and verify generated SQL before execution, improving data governance.
- Optimized Lambda cold start performance and execution memory allocation, reducing request latency.
- Applied modular architecture and Git-based version control workflows within Agile sprint cycles.

Software Engineer | Data

Feb 2023 – March 2024

- Designed a serverless backend architecture using AWS Lambda and S3 to automate secure medical report processing.
- Implemented event-driven orchestration using S3 triggers and asynchronous Lambda chaining.
- Refactored execution logic and optimized resource allocation, reducing processing time by 40%.
- Introduced structured logging and fault-tolerant retry mechanisms to improve reliability.
- Applied secure storage policies and IAM role-based access control to ensure HIPAA-aligned data handling.

LIS Nepal (Yomari)

KATHMANDU, NEPAL

Software Engineer | Data

March 2022 – Feb 2023

- Designed and automated Snowflake ELT pipelines for structured data migration.
- Implemented continuous ingestion using Snowpipe and batch execution via SnowSQL.
- Optimized complex SQL transformations to improve warehouse performance and query efficiency.
- Developed validation and reconciliation logic to detect inconsistencies across migrated datasets.
- Tuned warehouse sizing and clustering strategies to optimize cost-performance balance.

EDUCATION

CHARLES DARWIN UNIVERSITY

SYDNEY, NSW

Master of Data Science

March 2024 - November 2025

NOTABLE PROJECTS

Charles Darwin University | Client: Robust Movie Recommendation System

SYDNEY, NSW

Software Engineer

March 2025 – June 2025

- Built a hybrid recommender system combining **SVD (matrix factorization)** and **content-based similarity** for personalized movie recommendations.
- Developed and deployed an interactive **Streamlit** web application with real-time results.
- Integrated external metadata via TMDB API, enabling enriched models and higher recommendation accuracy.
- Designed modular pipelines to retrain models and update similarity matrices for continuous system improvement.

Institute of Engineering | Client: OCR & Text-to-Speech Automation

LALITPUR, NEPAL

Software Engineer

Dec 2022 – June 2023

- Developed a Python-based **OCR + TTS** system to automate text extraction and audio output.
- Improved processing speed by 25% through optimized preprocessing and model integration.
- Delivered a production-ready workflow for automated document-to-speech transformation.

CERTIFICATIONS & AWARDS

- Artificial Intelligence Fellowship – Fusemachines (2022-23)
- Data Fellowship – Code For Nepal (2023)
- Associate Member – Australian Computer Society (2024-Present)
- IT Head & Advisor – Leo Club of Kathmandu Budhi Gandaki (2020-24)
- Winner – Yomari Coding Competition (2022) | Datathon OpenIMIS-DRG Thailand (2022)
- Organiser – Locus Coding Competition (2023) | Presenter – AeHIN Hour (2022)

REFERENCE & ADDITIONAL INFORMATION

- **Reference:** Available upon request
- **Work Right:** Full-time, unrestricted working rights in Australia
- **Availability:** Immediately available for full-time employment
- **Flexibility:** Open to hybrid, remote or on-site roles
- **Relocation:** Willing to relocate within Australia if required